
1. Imports and Setup

We import all necessary libraries:

- For mathematical operations (`math`, `numpy`),
- For data handling (`csv`, `os`, `random`),
- For plotting (`matplotlib`),
- For building deep learning models (`tensorflow`, `tensorflow_addons`),
- For working with Recurrent Neural Networks (RNNs), LSTM cells, attention mechanisms, and dense layers.

```
%matplotlib inline

import math

import numpy as np

import os

import random

import tensorflow as tf

from matplotlib import pylab

from collections import Counter

import csv

# Seq2Seq Items

from tensorflow.python.ops.rnn_cell import LSTMCell

from tensorflow.python.ops.rnn_cell import MultiRNNCell

import tensorflow_addons.seq2seq as seq2seq

import tensorflow_addons.seq2seq.attention_wrapper as attention_wrapper

from tensorflow.python.layers.core import Dense
```

2. Defining Model Hyperparameters

We set the main hyperparameters of our model:

- Vocabulary size for both source (German) and target (English) languages,
- Number of units (hidden size) in the LSTM cells,
- Input embedding size,
- Batch size for training,
- Sequence lengths for encoder and decoder,
- Decoder type (basic or attention-based),
- Number of sentence pairs to read for training.

```
vocab_size= 50000

num_units = 128

input_size = 128

batch_size = 16

source_sequence_length=40

target_sequence_length=60

decoder_type = 'basic' # could be basic or attention

sentences_to_read = 50000
```

3. Loading Vocabulary Files

We load the vocabularies for both German (source) and English (target) languages.

- Each word is mapped to a unique integer index.
- Reverse dictionaries (index to word) are also created for easy decoding and visualization later.

```
src_dictionary = dict()

with open('vocab.50K.de.txt', encoding='utf-8') as f:

    for line in f:

        #we are discarding last char as it is new line char

        src_dictionary[line[:-1]] = len(src_dictionary)

src_reverse_dictionary =
dict(zip(src_dictionary.values(),src_dictionary.keys()))

print('Source')

print('\t',list(src_dictionary.items())[:10])

print('\t',list(src_reverse_dictionary.items())[:10])

print('\t','Vocabulary size: ', len(src_dictionary))

tgt_dictionary = dict()

with open('vocab.50K.en.txt', encoding='utf-8') as f:

    for line in f:

        #we are discarding last char as it is new line char

        tgt_dictionary[line[:-1]] = len(tgt_dictionary)

tgt_reverse_dictionary =
dict(zip(tgt_dictionary.values(),tgt_dictionary.keys()))
```

```
print('Target')

print('\t',list(tgt_dictionary.items())[:10])

print('\t',list(tgt_reverse_dictionary.items())[:10])

print('\t','Vocabulary size: ', len(tgt_dictionary))
```

4. Loading Training Sentences

We load the German-English training sentence pairs from files.

- Skip the first 50 lines to avoid noisy translations.
- Load exactly 50,000 sentences for both source and target datasets.
- Ensure the source and target datasets are perfectly aligned (same number of sentences).

```
source_sent = []

target_sent = []

test_source_sent = []

test_target_sent = []

with open('train.de', encoding='utf-8') as f:

    for l_i, line in enumerate(f):

        # discarding first 20 translations as there was some

        # english to english translations found in the first few. which

are wrong

        if l_i<50:
```

```
        continue

    source_sent.append(line)

    if len(source_sent)>=sentences_to_read:

        break

with open('train.en', encoding='utf-8') as f:

    for l_i, line in enumerate(f):

        if l_i<50:

            continue

        target_sent.append(line)

        if len(target_sent)>=sentences_to_read:

            break

assert len(source_sent)==len(target_sent), 'Source: %d, Target: %d'%(len(source_sent),len(target_sent))

print('Sample translations (%d)'%len(source_sent))

for i in range(0,sentences_to_read,10000):

    print('(',i,') DE: ', source_sent[i])

    print('(',i,') EN: ', target_sent[i])
```

5. Analyzing Sentence Statistics

We preprocess each sentence by:

- Separating punctuation from words,
- Replacing unknown words with a special `<unk>` token.

We also calculate:

- The average length and standard deviation of sentence lengths for both source and target datasets.
- This helps understand sentence distributions and set padding/truncation limits.

```
def split_to_tokens(sent, is_source):  
  
    #sent = sent.replace('-', ' ')  
  
    sent = sent.replace(',', ' ,')  
  
    sent = sent.replace('.', ' .')  
  
    sent = sent.replace('\n', ' ')  
  
  
    sent_toks = sent.split(' ')  
  
    for t_i, tok in enumerate(sent_toks):  
  
        if is_source:  
  
            if tok not in src_dictionary.keys():  
  
                sent_toks[t_i] = '<unk>'  
  
        else:  
  
            if tok not in tgt_dictionary.keys():  
  
                sent_toks[t_i] = '<unk>'  
  
    return sent_toks
```

```
# Let us first look at some statistics of the sentences

source_len = []

source_mean, source_std = 0,0

for sent in source_sent:

    source_len.append(len(split_to_tokens(sent,True)))

print('(Source) Sentence mean length: ', np.mean(source_len))

print('(Source) Sentence stddev length: ', np.std(source_len))


target_len = []

target_mean, target_std = 0,0

for sent in target_sent:

    target_len.append(len(split_to_tokens(sent,False)))

print('(Target) Sentence mean length: ', np.mean(target_len))

print('(Target) Sentence stddev length: ', np.std(target_len))
```

6. Preparing Training Inputs and Outputs

We prepare the final dataset for training:

- Add special tokens like `<s>` (start) and `</s>` (end).
- Reverse source sentences to help the model learn better (common trick in Seq2Seq).
- Pad all sentences to fixed maximum lengths to enable efficient batch processing.
- Create input length arrays to keep track of the actual (unpadded) lengths of sentences.

```

train_inputs = []

train_outputs = []

train_inp_lengths = []

train_out_lengths = []


max_tgt_sent_lengths = 0


src_max_sent_length = 41

tgt_max_sent_length = 61

for s_i, (src_sent, tgt_sent) in
enumerate(zip(source_sent,target_sent)):


    src_sent_tokens = split_to_tokens(src_sent,True)

    tgt_sent_tokens = split_to_tokens(tgt_sent,False)


    num_src_sent = []

    for tok in src_sent_tokens:

        num_src_sent.append(src_dictionary[tok])


    num_src_set = num_src_sent[::-1] # we reverse the source sentence.
This improves performance

    num_src_sent.insert(0,src_dictionary['<s>'])

train_inp_lengths.append(min(len(num_src_sent)+1,src_max_sent_length))


    # append until the sentence reaches max length

```



```

    if len(num_src_sent)<src_max_sent_length:

        num_src_sent.extend([src_dictionary['</s>'] for _ in
range(src_max_sent_length - len(num_src_sent))])

        # if more than max length, truncate the sentence

    elif len(num_src_sent)>src_max_sent_length:

        num_src_sent = num_src_sent[:src_max_sent_length]

    assert len(num_src_sent)==src_max_sent_length,len(num_src_sent)

train_inputs.append(num_src_sent)

num_tgt_sent = [tgt_dictionary['</s>']]

for tok in tgt_sent_tokens:

    num_tgt_sent.append(tgt_dictionary[tok])

train_out_lengths.append(min(len(num_tgt_sent)+1,tgt_max_sent_length))

    if len(num_tgt_sent)<tgt_max_sent_length:

        num_tgt_sent.extend([tgt_dictionary['</s>'] for _ in
range(tgt_max_sent_length - len(num_tgt_sent))])

    elif len(num_tgt_sent)>tgt_max_sent_length:

        num_tgt_sent = num_tgt_sent[:tgt_max_sent_length]

train_outputs.append(num_tgt_sent)

    assert len(train_outputs[s_i])==tgt_max_sent_length, 'Sent length
needs to be 60, but is %d'%len(binned_outputs[s_i])

```

```

assert len(train_inputs) == len(source_sent),\

    'Size of total bin elements: %d, Total sentences: %d'\

    %(len(train_inputs),len(source_sent))

print('Max sent lengths: ', max_tgt_sent_lengths)

train_inputs = np.array(train_inputs, dtype=np.int32)

train_outputs = np.array(train_outputs, dtype=np.int32)

train_inp_lengths = np.array(train_inp_lengths, dtype=np.int32)

train_out_lengths = np.array(train_out_lengths, dtype=np.int32)

print('Samples from bin')

print('\t', [src_reverse_dictionary[w] for w in
train_inputs[0,:].tolist()])

print('\t', [tgt_reverse_dictionary[w] for w in
train_outputs[0,:].tolist()])

print('\t', [src_reverse_dictionary[w] for w in
train_inputs[10,:].tolist()])

print('\t', [tgt_reverse_dictionary[w] for w in
train_outputs[10,:].tolist()])

print()

print('\tSentences ', train_inputs.shape[0])

```

7. Data Batching Class

We define a custom batch generator class.

- This class handles how batches are created for training.
- It efficiently unrolls the source and target sentences into sequential batches.
- It manages sentence cursors and resets after reaching the end of sentences.
- Word embeddings for German and English are loaded from pre-trained `.npy` files.

```
input_size = 128

class DataGeneratorMT(object):

    def __init__(self, batch_size, num_unroll, is_source):

        self._batch_size = batch_size

        self._num_unroll = num_unroll

        self._cursor = [0 for offset in range(self._batch_size)]

        self._src_word_embeddings = np.load('de-embeddings.npy')

        self._tgt_word_embeddings = np.load('en-embeddings.npy')

        self._sent_ids = None

        self._is_source = is_source

    def next_batch(self, sent_ids, first_set):
```

```

if self._is_source:

    max_sent_length = src_max_sent_length

else:

    max_sent_length = tgt_max_sent_length

batch_labels_ind = []

batch_data = np.zeros((self._batch_size),dtype=np.float32)

batch_labels = np.zeros((self._batch_size),dtype=np.float32)


for b in range(self._batch_size):


    sent_id = sent_ids[b]


    if self._is_source:

        sent_text = train_inputs[sent_id]


        batch_data[b] = sent_text[self._cursor[b]]

        batch_labels[b]=sent_text[self._cursor[b]+1]


    else:

        sent_text = train_outputs[sent_id]


        batch_data[b] = sent_text[self._cursor[b]]

        batch_labels[b] = sent_text[self._cursor[b]+1]


    self._cursor[b] = (self._cursor[b]+1)%(max_sent_length-1)

```

```
        return batch_data, batch_labels

def unroll_batches(self, sent_ids):

    if sent_ids is not None:

        self._sent_ids = sent_ids

        self._cursor = [0 for _ in range(self._batch_size)]

    unroll_data, unroll_labels = [], []

    inp_lengths = None

    for ui in range(self._num_unroll):

        data, labels = self.next_batch(self._sent_ids, False)

        unroll_data.append(data)

        unroll_labels.append(labels)

        inp_lengths = train_inp_lengths[sent_ids]

    return unroll_data, unroll_labels, self._sent_ids, inp_lengths

def reset_indices(self):

    self._cursor = [0 for offset in range(self._batch_size)]
```

```

# Running a tiny set to see if the implementation correct

dg = DataGeneratorMT(batch_size=5,num_unroll=40,is_source=True)

u_data, u_labels, _, _ = dg.unroll_batches([0,1,2,3,4])

print('Source data')

for _, lbl in zip(u_data,u_labels):

    print([src_reverse_dictionary[w] for w in lbl.tolist()])


# Running a tiny set to see if the implementation correct

dg = DataGeneratorMT(batch_size=5,num_unroll=60,is_source=False)

u_data, u_labels, _, _ = dg.unroll_batches([0,2,3,4,5])

print('\nTarget data batch (first time)')

for d_i,(_, lbl) in enumerate(zip(u_data,u_labels)):

    #if d_i>5 and d_i < 35:

    #    continue

    print([tgt_reverse_dictionary[w] for w in lbl.tolist()])

print('\nTarget data batch (non-first time)')

u_data, u_labels, _, _ = dg.unroll_batches(None)

for d_i,(_, lbl) in enumerate(zip(u_data,u_labels)):

    #if d_i>5 and d_i < 35:

    #    continue

```

```
print([tgt_reverse_dictionary[w] for w in lbl.tolist()])
```

8. Input Placeholders, Embeddings, and Masks

We set up TensorFlow placeholders for:

- Encoder and decoder inputs,
- Decoder output labels,
- Masks for padding during loss computation.

Pre-trained word embeddings are used instead of training embeddings from scratch, saving time and improving translation quality.

```
tf.reset_default_graph()

enc_train_inputs = []
dec_train_inputs = []

# Need to use pre-trained word embeddings
encoder_emb_layer = tf.convert_to_tensor(np.load('de-embeddings.npy'))
decoder_emb_layer = tf.convert_to_tensor(np.load('en-embeddings.npy'))

# Defining unrolled training inputs
for ui in range(source_sequence_length):

    enc_train_inputs.append(tf.placeholder(tf.int32,
        shape=[batch_size], name='enc_train_inputs_%d'%ui))
```

```

dec_train_labels=[]

dec_label_masks = []

for ui in range(target_sequence_length):

    dec_train_inputs.append(tf.placeholder(tf.int32,
shape=[batch_size],name='dec_train_inputs_%d'%ui))

    dec_train_labels.append(tf.placeholder(tf.int32,
shape=[batch_size],name='dec-train_outputs_%d'%ui))

    dec_label_masks.append(tf.placeholder(tf.float32,
shape=[batch_size],name='dec-label_masks_%d'%ui))

encoder_emb_inp = [tf.nn.embedding_lookup(encoder_emb_layer, src) for
src in enc_train_inputs]

encoder_emb_inp = tf.stack(encoder_emb_inp)

decoder_emb_inp = [tf.nn.embedding_lookup(decoder_emb_layer, src) for
src in dec_train_inputs]

decoder_emb_inp = tf.stack(decoder_emb_inp)

enc_train_inp_lengths = tf.placeholder(tf.int32,
shape=[batch_size],name='train_input_lengths')

dec_train_inp_lengths = tf.placeholder(tf.int32,
shape=[batch_size],name='train_output_lengths')

```

9. Encoder Architecture

We define the encoder as a basic LSTM network:

- It processes the input German sentences,
- Produces hidden states (`encoder_outputs`) and the final context (`encoder_state`),
- The encoder is dynamic, meaning it can handle variable sentence lengths.

```
encoder_cell = tf.nn.rnn_cell.BasicLSTMCell(num_units)

initial_state = encoder_cell.zero_state(batch_size, dtype=tf.float32)

encoder_outputs, encoder_state = tf.nn.dynamic_rnn(
    encoder_cell, encoder_emb_inp, initial_state=initial_state,
    sequence_length=enc_train_inp_lengths,
    time_major=True, swap_memory=True)
```

10. Decoder Architecture

We define the decoder as another LSTM network:

- It uses the final encoder state to start decoding,
- It can be either a basic decoder or attention-based depending on the setting,
- A dense projection layer maps LSTM outputs to vocabulary logits (prediction probabilities).

```
# Build RNN cell

decoder_cell = tf.nn.rnn_cell.BasicLSTMCell(num_units)

projection_layer = Dense(units=vocab_size, use_bias=True)
```

```

# Helper
helper = tf.contrib.seq2seq.TrainingHelper(
    decoder_emb_inp, [tgt_max_sent_length-1 for _ in range(batch_size)],
    time_major=True)

# Decoder
if decoder_type == 'basic':
    decoder = tf.contrib.seq2seq.BasicDecoder(
        decoder_cell, helper, encoder_state,
        output_layer=projection_layer)

elif decoder_type == 'attention':
    decoder = tf.contrib.seq2seq.BahdanauAttention(
        decoder_cell, helper, encoder_state,
        output_layer=projection_layer)

# Dynamic decoding
outputs, _, _ = tf.contrib.seq2seq.dynamic_decode(
    decoder, output_time_major=True,
    swap_memory=True
)

```

11. Loss Function and Optimization

We define the loss:

- Use cross-entropy loss between predicted and actual words.
- Apply masking so that padding tokens don't contribute to loss.

We also define two optimizers:

- **Adam** optimizer is used for faster convergence in early training.
- **SGD** optimizer is used later for fine-tuning.
- Gradient clipping is applied to prevent exploding gradients.

```
logits = outputs.rnn_output

crossent = tf.nn.sparse_softmax_cross_entropy_with_logits(
    labels=dec_train_labels, logits=logits)

loss = (tf.reduce_sum(crossent*tf.stack(dec_label_masks)) /
        (batch_size*target_sequence_length))

train_prediction = outputs.sample_id
```

12. Training Loop

We run the training for 10001 steps:

- In each step, sample random sentence pairs,
- Pass them through encoder and decoder,
- Compute loss and optimize the model,

- Switch optimizer after 10000 steps for better stability,
- Every 250 steps, print actual vs predicted translations to monitor progress,
- Every 500 steps, log and reset the running average of the loss.

```
print('Defining Optimizer')

# Adam Optimizer. And gradient clipping.

global_step = tf.Variable(0, trainable=False)

inc_gstep = tf.assign(global_step, global_step + 1)

learning_rate = tf.train.exponential_decay(

    0.01, global_step, decay_steps=10, decay_rate=0.9, staircase=True)

with tf.variable_scope('Adam'):

    adam_optimizer = tf.train.AdamOptimizer(learning_rate)

    adam_gradients, v = zip(*adam_optimizer.compute_gradients(loss))

    adam_gradients, _ = tf.clip_by_global_norm(adam_gradients, 25.0)

    adam_optimize = adam_optimizer.apply_gradients(zip(adam_gradients, v))

with tf.variable_scope('SGD'):

    sgd_optimizer = tf.train.GradientDescentOptimizer(learning_rate)

    sgd_gradients, v = zip(*sgd_optimizer.compute_gradients(loss))

    sgd_gradients, _ = tf.clip_by_global_norm(sgd_gradients, 25.0)

    sgd_optimize = sgd_optimizer.apply_gradients(zip(sgd_gradients, v))

sess = tf.InteractiveSession()
```

```
if not os.path.exists('logs'):

    os.mkdir('logs')

log_dir = 'logs'

bleu_scores_over_time = []

loss_over_time = []

tf.global_variables_initializer().run()


src_word_embeddings = np.load('de-embeddings.npy')

tgt_word_embeddings = np.load('en-embeddings.npy')


# Defining data generators

enc_data_generator =
DataGeneratorMT(batch_size=batch_size,num_unroll=source_sequence_length
,is_source=True)

dec_data_generator =
DataGeneratorMT(batch_size=batch_size,num_unroll=target_sequence_length
,is_source=False)

num_steps = 10001

avg_loss = 0

bleu_labels, bleu_preds = [],[]

print('Started Training')

for step in range(num_steps):
```

```

# input_sizes for each bin: [40]

# output_sizes for each bin: [60]

print('.',end='')

if (step+1)%100==0:

    print('')

    sent_ids =
np.random.randint(low=0,high=train_inputs.shape[0],size=(batch_size))

    # ===== ENCODER DATA COLLECTION
    =====

    eu_data, eu_labels, _, eu_lengths =
enc_data_generator.unroll_batches(sent_ids=sent_ids)

    feed_dict = {}

    feed_dict[enc_train_inp_lengths] = eu_lengths

    for ui,(dat,lbl) in enumerate(zip(eu_data,eu_labels)):

        feed_dict[enc_train_inputs[ui]] = dat

    # ===== DECODER DATA COLLECITON
    =====

    # First step we change the ids in a batch

    du_data, du_labels, _, du_lengths =
dec_data_generator.unroll_batches(sent_ids=sent_ids)

    feed_dict[dec_train_inp_lengths] = du_lengths

```

```

for ui,(dat,lbl) in enumerate(zip(du_data,du_labels)):

    feed_dict[dec_train_inputs[ui]] = dat

    feed_dict[dec_train_labels[ui]] = lbl

    feed_dict[dec_label_masks[ui]] = (np.array([ui for _ in
range(batch_size)])<du_lengths).astype(np.int32)

# ===== OPTIMIZATION =====

if step < 10000:

    _,l,tr_pred = sess.run([adam_optimize,loss,train_prediction],
feed_dict=feed_dict)

else:

    _,l,tr_pred = sess.run([sgd_optimize,loss,train_prediction],
feed_dict=feed_dict)

tr_pred = tr_pred.flatten()


if (step+1)%250==0:

    print('Step ',step+1)

    print_str = 'Actual: '

    for w in
np.concatenate(du_labels,axis=0)[:batch_size].tolist():

        print_str += tgt_reverse_dictionary[w] + ' '

        if tgt_reverse_dictionary[w] == '</s>':

            break

```

```

print(print_str)

print()

print_str = 'Predicted: '

for w in tr_pred[::batch_size].tolist():

    print_str += tgt_reverse_dictionary[w] + ' '

    if tgt_reverse_dictionary[w] == '</s>':

        break

print(print_str)

print('\n')

rand_idx = np.random.randint(low=1,high=batch_size)

print_str = 'Actual: '

for w in
np.concatenate(du_labels,axis=0)[rand_idx::batch_size].tolist():

    print_str += tgt_reverse_dictionary[w] + ' '

    if tgt_reverse_dictionary[w] == '</s>':

        break

print(print_str)

print()

print_str = 'Predicted: '

for w in tr_pred[rand_idx::batch_size].tolist():

```



```

        print_str += tgt_reverse_dictionary[w] + ' '

        if tgt_reverse_dictionary[w] == '</s>':

            break

    print(print_str)

    print()

    avg_loss += 1

#sess.run(reset_train_state) # resetting hidden state for each batch

if (step+1)%500==0:

    print('==== Step ', str(step+1), ' =====')

    print('\t Loss: ', avg_loss/500.0)

    loss_over_time.append(avg_loss/500.0)

    avg_loss = 0.0

    sess.run(inc_gstep)

```

Conclusion

This project trains a basic **Neural Machine Translation (NMT)** model that translates sentences from German to English using:

- LSTM-based Encoder-Decoder architecture,
- Pre-trained word embeddings,
- Proper batching, masking, and optimization techniques,
- Dynamic handling of variable sentence lengths.